

# Kubernetes

**Kubernetes** is a popular open-source platform for automating deployment, scaling, and management of containerized applications.

Here are some important concepts to know when working with Kubernetes:

- **Nodes:**
  1. A node is a worker machine in the Kubernetes cluster that runs one or more pods.
  2. Nodes are responsible for running the applications and services that make up your containerized application.
  3. Each node has its own set of resources, such as CPU and memory, that are used by the pods running on that node.
  4. Nodes can be physical machines or virtual machines, and are managed by the cluster's control plane.
- **Pods:**
  1. A pod is the smallest deployable unit in Kubernetes.
  2. It is a logical host for one or more containers, which share the same network namespace and can communicate with each other using localhost.
- **ReplicaSets:**
  1. A ReplicaSet ensures that a specified number of identical pods are running at any given time.
  2. If a pod fails or is terminated, the ReplicaSet will automatically create a new pod to maintain the desired number of replicas.
- **Deployments:**
  1. A Deployment manages the creation and scaling of pods.
  2. It allows you to update the desired state of your application without downtime, by creating a new pod with the updated configuration and gradually scaling it up while scaling down the old pod.
- **Services:**

1. A Service provides a stable IP address and DNS name for a set of pods, allowing other parts of your application to access them by a consistent name.
2. Services can be exposed to the external network or only within the cluster.
3. There are four types of services that can be used to expose your application to other parts of the cluster or to the outside world:
  - I. **ClusterIP:** A ClusterIP service provides a stable IP address and DNS name for a set of pods within the cluster. This type of service is only accessible from within the cluster and is often used for internal communication between components of an application.
  - II. **NodePort:** A NodePort service exposes a port on every node in the cluster, allowing external traffic to be routed to the service. This type of service is useful for testing and development, but it can be difficult to manage at scale and is often not suitable for production workloads.
  - III. **LoadBalancer:** A LoadBalancer service creates an external load balancer in cloud environments such as AWS or GCP, allowing incoming traffic to be distributed across the pods in the service. This type of service is often used for production workloads that require high availability and scalability.
  - IV. **ExternalName:** An ExternalName service provides an alias for an external service outside the cluster. This type of service is useful for integrating with external services, such as databases or APIs, that have a stable DNS name or IP address.

By choosing the right type of service for your application, you can ensure that it is accessible and scalable in a way that meets your needs. In addition, Kubernetes supports advanced networking features such as Ingress and Network Policies, which can be used to further customize and secure your application's network traffic.

- **ConfigMaps and Secrets:**

1. ConfigMaps allow you to store configuration data that can be accessed by your application's containers, while Secrets allow you to store sensitive information such as passwords and API keys.
2. Both can be mounted as files or environment variables in your application's containers.

To interact with Kubernetes, we can use kubectl, a command-line tool for deploying and managing applications on Kubernetes clusters. Here are some common kubectl commands:

- **kubectl create:** This command creates a new resource in the Kubernetes cluster. You can specify the resource type and configuration details in a YAML file or on the command line.

**Example:** `kubectl create deployment nginx --image=nginx:latest` creates a deployment resource named "nginx" with the latest nginx image.

- **kubectl apply:** This command applies changes to a resource by updating or creating it. It can be used to apply changes to multiple resources at once by specifying a directory containing YAML configuration files.

**Example:** `kubectl apply -f nginx.yaml` applies changes specified in the "nginx.yaml" file to the Kubernetes cluster.

- **kubectl get:** This command retrieves information about resources in the Kubernetes cluster. You can filter the results by specifying a resource type and/or a label selector.

**Example:** `kubectl get pods` retrieves information about all pods in the cluster.

- **kubectl describe:** This command provides detailed information about a specific resource in the Kubernetes cluster, such as its current state and associated events.

**Example:** `kubectl describe pod nginx-123` provides detailed information about the "nginx-123" pod.

- **kubectl logs:** This command retrieves the logs of a specific container within a pod. You can specify the container name and/or the number of lines to retrieve.

**Example:** `kubectl logs nginx-123 -c nginx` retrieves the logs of the "nginx" container within the "nginx-123" pod.

- **kubectl exec:** This command executes a command within a running container. You can specify the container name and/or the command to execute.

**Example:** `kubectl exec nginx-123 -c nginx -- ls -la` executes the "ls -la" command within the "nginx" container within the "nginx-123" pod.

- **kubectl delete:** This command deletes a resource from the Kubernetes cluster. You can specify the resource type and/or a label selector to delete multiple resources at once.

Example: `kubectl delete pod nginx-123` deletes the "nginx-123" pod.